

Spark with Python and Scala Programming

Hieu Nguyen

March 20, 2017

Outline

- 1 SIS BigData platform
- 2 Apache Spark
- 3 PySpark: spark python big data processing
- 4 Scala: spark scala big data processing
- 5 Spark SQL and DataFrames

SIS BigData platform

- Our SIS BigData cluster (5 nodes) is now ready to use
- **bigdata0** (coordinator node):
8 cores CPU, 23.5 GiB Mem, and 88 GiB Disk
- **bigdata1, bigdata2:**
24 cores CPU, 78.6 GiB Mem, 7.5 TiB Disk
- **bigdata3, bigdata4, bigdata5:**
24 cores CPU, 62.9 GiB Mem, 4.2 TiB Disk
- The version of Cloudera is CDH 5.10.0
 - 4 services was deployed: HDFS, Hadoop MapReduce 2 (with YARN), Spark, and HBase
 - other services are available: Accumulo, Flume, Hive, Hue, Impala, Isilon, Kafka, Oozie, S3 Connector, Sentry, Solr, and Sqoo

How to Connect to SIS BigData platform

- You need an account on BigData cluster
- **Outside of University network:**
 - log into `shell.sis.uta.fi` (with your basic account)
(if you are using Windows OS, simply use Tectia-SSH Terminal)
 - Then, `ssh <bigdata-account>@<server>.sis.uta.fi`
- **Inside of University network:**
 - `ssh <bigdata-account>@<server>.sis.uta.fi`

HDFS commands

- `hadoop fs -mkdir /user/<user-name>/foo`
 - creates a directory called "foo" in the HDFS home directory of the user <user-name>
- `hadoop fs -rm -r /user/<user-name>/foo`
 - removes a directory called "foo" in the HDFS home directory of the user <user-name>
- `hadoop fs -put foo.txt /user/<user-name>/foo/`
 - copy a file "foo.txt" to the HDFS "foo" directory
 - similar: `hadoop fs -copyFromLocal foo.txt /user/<user-name>/foo/`
- `hadoop fs -cat /user/<user-name>/foo/foo.txt`
 - sees the content of the file "foo.txt" in the HDFS
- `hadoop fs -rm /user/<user-name>/foo/foo.txt`
 - removes a file called "foo.txt" in the HDFS
- `hadoop fs -ls /user/<user-name>/foo`
 - does a directory listing for directory "foo" in the HDFS

HDFS commands (cont.)

- `hadoop fs -get /user/<user-name>/bar`
 - copy the directory "bar" in the user HDFS directory to the local file system
 - similar: `hadoop fs -copyToLocal /user/<user-name>/bar ./`
- `hadoop fs -getmerge /user/<user-name>/bar >> output_merge.txt`
 - does a file merge for the output if Spark job creates multiple output files, then stores the output in the local file system "output_merge.txt"

Apache Spark: lightning-fast cluster computing

- a fast and general engine for **large-scale data processing**
- **speed**
 - up to 100x faster than Hadoop MapReduce in memory
 - or 10x faster on disk
- **provides high-level**
 - APIs in: **Java, Scala, Python, and R**
 - tools including:
 - **Spark SQL and DataFrames** for SQL and structured data processing
 - **MLlib** for machine learning
 - **GraphX** for graph processing
 - **Spark Streaming** for stream processing of live datastreams
- **run everywhere**: standalone cluster mode, on EC2, on Hadoop YARN, or on Apache Mesos
- **access diverse data sources**: including HDFS, Cassandra, HBase, Hive, and any Hadoop data source

Apache Spark: components

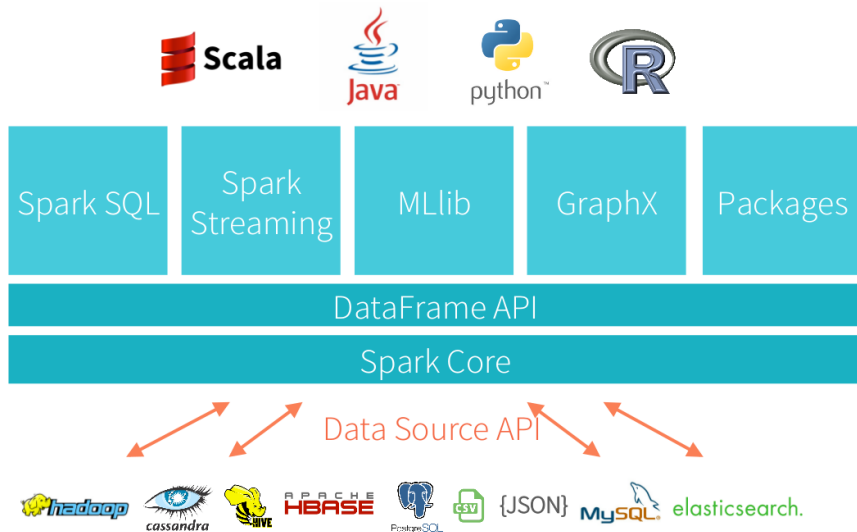


figure from Databricks

Apache Spark: open source ecosystem

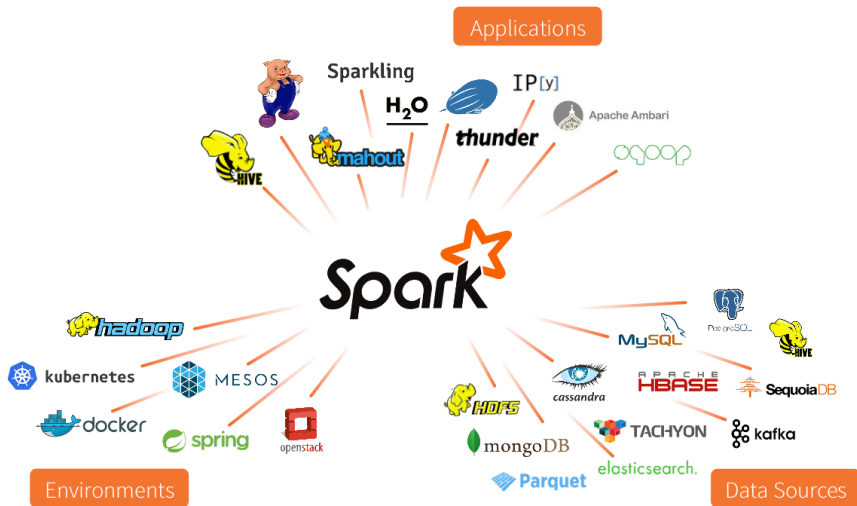


figure from Databricks

Apache Spark: resilient distributed datasets (RDD)

- **collections of objects** spread across a cluster
- stored in **RAM** or on **Disk**
- built through **parallel transformations**
- **automatically** rebuilt on failure
- two types of operations on RDDs: **transformations** and **actions**
- you can **persist (cache)** an RDD
 - by using: `persist()` or `cache()`
 - each node stores any partitions of it
 - computes in memory
 - reuses them in other actions
 - can choose storage level:
MEMORY_ONLY, DISK_ONLY, MEMORY_AND_DISK
 - release by: `unpersist()`

Apache Spark: transformations and actions

- **Spark transformations: are lazy (not executed until an action follows)**


- `map()`
- `flatMap()`
- `reduceByKey()`
- `filter()`
- `sample()`
- `union()`
- `intersection()`
- `distinct()`
- `groupByKey()`
- `sortByKey()`
- `join()`
- `cogroup`

- **Spark actions: time consuming (execution of an action results in all the previously created transformation)**

- `reduce()`
- `collect()`
- `count()`
- `countByValue()`
- `first()`
- `take()`
- `takeSample()`
- `foreach()`
- `saveAsTextFile()`
- `saveAsHadoop()`

Apache Spark: transformations and actions

- **transformations:** just "remember" the operation to be performed, the dataset to which the operation is to be performed
- **action:** brings back the data from the RDD to the local machine



```
// transformed RDDs
val errors = lines.filter(_.startsWith("ERROR"))
val messages = errors.map(_.split("\t")).map(r => r(1))
messages.cache()
```

```
// action 1
messages.filter(_.contains("mysql")).count()
```

figure from Databricks

Apache Spark: SparkContext

- **main entry point** to Spark functionality
- specifies running environment and app name
- uses to create RDDs from many input sources
- creates counters and accumulators
- available in shell (interactive mode) as variable `sc`
- in your program: you'd make your own

PySpark WordCount: including spark libraries

```
from pyspark import SparkContext, SparkConf
from pyspark.sql import SQLContext
from pyspark.sql.types import *
from pyspark.sql import Row

import pandas as pd
import matplotlib.pyplot as plt
import re
```

PySpark WordCount: main function

```
if __name__ == '__main__':  
    //Configure Spark  
    APP_NAME = "PySparkExample"  
    conf = SparkConf().setAppName(APP_NAME).setMaster("local[1]")  
    //conf = SparkConf().setAppName(APP_NAME).setMaster("local[*]")  
    //conf = SparkConf().setAppName(APP_NAME).setMaster("yarn-client")  
    //conf = SparkConf().setAppName(APP_NAME).setMaster("yarn-master")  
    sc = SparkContext(conf=conf)  
    //Transforms (load) the input data from HDFS into an RDD  
    rddData = sc.textFile("hdfs:///user/hieunguyen/input/wc.txt")  
    //Transformed RDD: process line by line, split by space  
    wcFM = rddData.flatMap(lambda line: line.split(" "))  
    //Transformed RDD: pass each element (key) by 1 (value)  
    wcM = wcFM.map(lambda word: (word, 1))  
    //Transformed RDD: merge key with an associative function  
    wcRBK = wcM.reduceByKey(lambda a,b: a+b)  
    //No data is read or processed until after this line  
    //Actioned RDD: save the results into HDFS  
    wcRBK.saveAsTextFile("hdfs:///user/hieunguyen/wcoutput")
```

PySpark WordCount: run your code with pyspark shell

- Interactive mode: pyspark

- `pyspark --master local[10] --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2`

- `pyspark --master local[*] --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2`

- `pyspark --master yarn-client --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2`

use additional package from third party

- `pyspark --master yarn-client --packages com.databricks:spark-csv_2.10:1.5.0`

use additional jar file

- `pyspark --master yarn-client --jars test.jar`

PySpark WordCount: run your code with spark-submit

- `spark-submit`
 - `spark-submit pysparkExample.py`
 - `spark-submit --master yarn-client --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2 pysparkExample.py`
 - `spark-submit --master yarn-cluster --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2 pysparkExample.py`
- run your code in background, separated spark output and error
 - `nohup spark-submit --master yarn-client --executor-memory 25G --driver-memory 20G --num-executors 12 --executor-cores 2 pysparkExample.py > output.out 2>error.err`

PySpark WordCount: how spark job run on a cluster

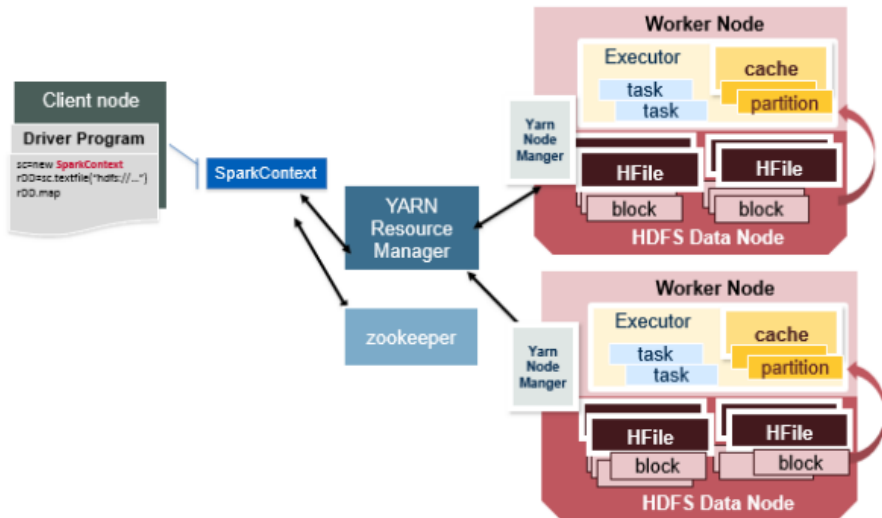


figure from www.mapr.com

PySpark WordCount: demo

- Demo PySpark WordCount

PySpark: accumulators and user defined functions

- **accumulators**: the global variable that can be shared across tasks
- **UDF**: simple way to add separate functions into Spark that can be used during various transformation stages

PySpark accumulators and UDF: demo

- Demo PySpark accumulators and UDF

Scala WordCount: including spark libraries

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkConf
import org.apache.spark.rdd.RDD
import java.io._
```

Scala WordCount: main function

```
object ScalaWordCount {  
  def main(args: Array[String]) {  
    val APP_NAME = "ScalaWordCount"  
  
    val conf = new  
      SparkConf().setAppName(APP_NAME).setMaster("yarn-client")  
  
    val sc = new SparkContext(conf)  
  
    val rddData =  
      sc.textFile("hdfs:///user/hieunguyen/input/wc.txt")  
  
    val wcFM = rddData.flatMap(line => line.split(" "))  
    val wcM = wcFM.map(word => (word, 1))  
    val wcRBK = wcM.reduceByKey((a, b) => a+b)  
    wcRBK.saveAsTextFile("hdfs:///user/hieunguyen/wcoutput")  
  }  
}
```

Scala WordCount: build source code by using sbt

- Your directory layout should look like this

```
$ find .  
./simple.sbt  
./src  
./src/main  
./src/main/scala  
./src/main/scala/ScalaWordCount.scala
```

- content of the simple.sbt file

```
name := "ScalaWordCount"  
version := "1.0"  
scalaVersion := "2.10.6"  
libraryDependencies ++= Seq(  
  "org.apache.spark" %% "spark-core" % "1.6.0"  
)
```

- build by: sbt package

Scala WordCount: run your Scala code

- Interactive mode: `spark-shell`
 - `spark-shell --master yarn-client`
`--executor-memory 25G --driver-memory 20G`
`--num-executors 12 --executor-cores 2`
- `spark-submit`
 - `spark-submit --class "ScalaWordCount" --master yarn-client`
`target/scala-2.10/scalawordcount_2.10-1.0.jar`
`"seminar/input-wordcount" 1`
 - `"seminar/input-wordcount"` is input file path on HDFS
 - `1` is threshold


Scala WordCount: demo

- Demo Scala WordCount

Spark SQL and DataFrames (DF)

- a DF is a **distributed collection of rows under named columns**
- DF in Spark has the **ability to handle petabytes of data**
- DF has a support for **wide range of data format and sources**
- has API support for different languages like **Python, R, Scala, Java**

Ways to Create DataFrame in Spark



Spark SQL and DataFrames: demo

- Demo Spark SQL and DataFrames

Spark: additional information

- turn an existing collection into an RDD
 - `sc.parallelize(<collection>)`
i.e., `sc.parallelize([1,2,3])`
- load complete content of the file at once
 - `sc.wholeTextFiles(<file/folder path>)`
- load the file with partitions
 - `sc.wholeTextFiles(<file/folder path>, <number>)`
 - `sc.textFile(<file/folder path>, <number>)`
 - `sc.parallelize(<collection>, <number>)`
- use existing Hadoop input form (only for Java and Scala)
 - `sc.hadoopFile(keyClass, valClass, inputForm, conf)`
- key-value pair in Python vs. Scala
 - Python: `pair = (a, b), pair[0] = a, pair[1] = b`
 - Scala: `pair = (a, b), pair._1 = a, pair._2 = b`



Programming Guide

<http://spark.apache.org/docs/latest/programming-guide.html> .